

Lovely Pandas



Contents

- Datasets
 - Data representation
 - Types of variables
- Pandas for Data Analysis
 - The pandas.Series data structure
 - The pandas.DataFrame data structure
 - Read data from files
- Basics of Descriptive Analytics
 - Descriptive measures
 - Visualizing data

Datasets

- Data representation

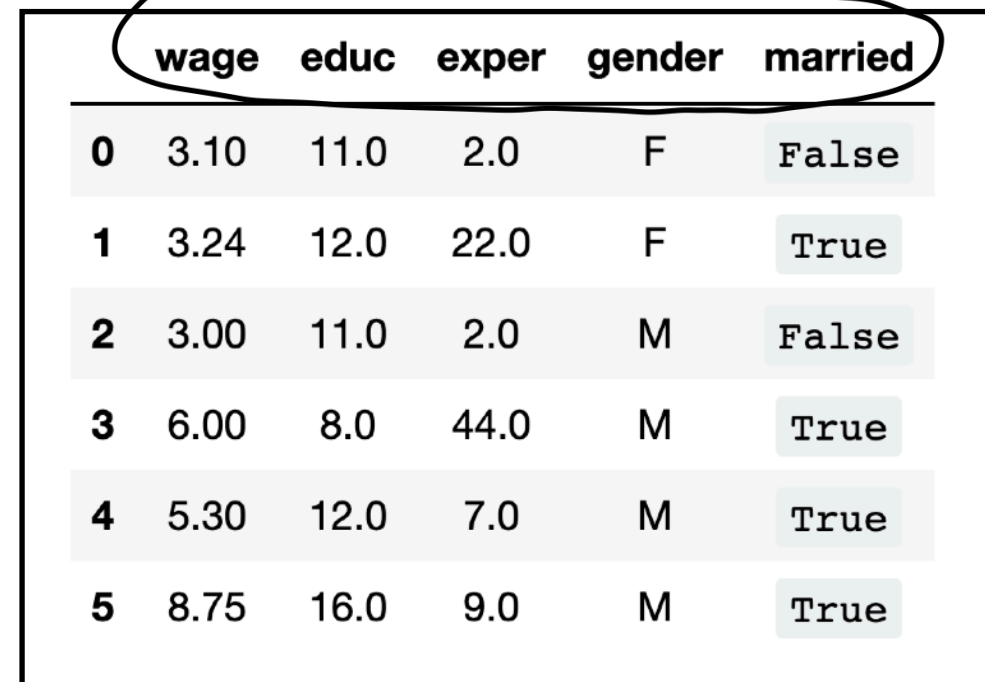
Example 1: The following table contains a cross-sectional dataset on a number of working individuals for the year 1976. Columns of the data table are summarized as follows.

- **wage:** average hourly earnings (in dollars)
- **educ:** years of education
- **exper:** years of potential experience
- **gender:** genders of these working individuals
- **married:** **True** if married, and **False** otherwise

Datasets

- Data representation
 - Variables (fields/attributes) as columns

Column labels



	wage	educ	exper	gender	married
0	3.10	11.0	2.0	F	False
1	3.24	12.0	22.0	F	True
2	3.00	11.0	2.0	M	False
3	6.00	8.0	44.0	M	True
4	5.30	12.0	7.0	M	True
5	8.75	16.0	9.0	M	True

Datasets

- Data representation
 - Variables (fields/attributes) as columns
 - Observations (cases/records) as rows

The diagram shows a table with 6 rows and 6 columns. The first column contains row indices from 0 to 5. The other five columns are labeled 'wage', 'educ', 'exper', 'gender', and 'married'. An arrow labeled 'Column labels' points to the header row. Another arrow labeled 'Row labels' points to the first column. The table is enclosed in a black border, and the header row is circled in black.

	wage	educ	exper	gender	married
0	3.10	11.0	2.0	F	False
1	3.24	12.0	22.0	F	True
2	3.00	11.0	2.0	M	False
3	6.00	8.0	44.0	M	True
4	5.30	12.0	7.0	M	True
5	8.75	16.0	9.0	M	True

Datasets

- Data representation
 - Types of variables
 - ✓ Numerical (quantitative) variables
 - ✓ Categorical (qualitative) variables

```
print(int(True))  
print(int(False))
```

1
0

The diagram shows a dataset table with 6 rows and 6 columns. The first three columns (index, wage, educ, exper) are grouped under the label 'Numerical variables' with an arrow. The last two columns (gender, married) are grouped under the label 'Categorical variables' with an arrow. The 'married' column contains boolean values (True/False).

	wage	educ	exper	gender	married
0	3.10	11.0	2.0	F	False
1	3.24	12.0	22.0	F	True
2	3.00	11.0	2.0	M	False
3	6.00	8.0	44.0	M	True
4	5.30	12.0	7.0	M	True
5	8.75	16.0	9.0	M	True

Pandas for Data Analysis

- Introduction to Pandas

- Labeled data
- Heterogenous data
- Possible missing values

- Import Pandas

```
import pandas as pd
```

Column labels

	wage	educ	exper	gender	married
0	3.10	11.0	2.0	F	False
1	3.24	12.0	22.0	F	True
2	3.00	11.0	2.0	M	False
3	6.00	8.0	44.0	M	True
4	5.30	12.0	7.0	M	True
5	8.75	16.0	9.0	M	True

Row labels

Pandas for Data Analysis

- The `pandas.Series` data structure
 - A one-dimensional array of indexed data

One-dimensional array

	wage	educ	exper	gender	married
0	3.10	11.0	2.0	F	False
1	3.24	12.0	22.0	F	True
2	3.00	11.0	2.0	M	False
3	6.00	8.0	44.0	M	True
4	5.30	12.0	7.0	M	True
5	8.75	16.0	9.0	M	True

wage
3.10
3.24
3.00
6.00
5.30
8.75

Length is six

Pandas for Data Analysis

- The `pandas.Series` data structure
 - A one-dimensional array of indexed data

One-dimensional array

1	3.24	12.0	22.0	F	True
---	------	------	------	---	------

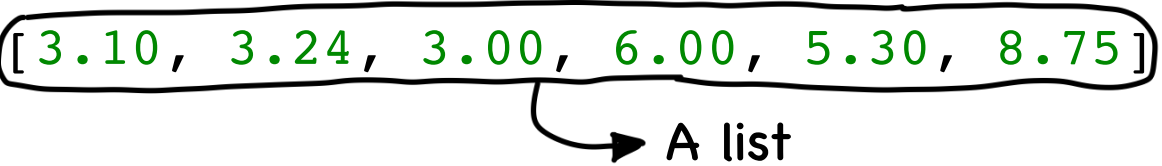
← Length is five →

	wage	educ	exper	gender	married
0	3.10	11.0	2.0	F	False
1	3.24	12.0	22.0	F	True
2	3.00	11.0	2.0	M	False
3	6.00	8.0	44.0	M	True
4	5.30	12.0	7.0	M	True
5	8.75	16.0	9.0	M	True

Pandas for Data Analysis

- The `pandas.Series` data structure
 - A one-dimensional array of indexed data
 - Create a `pandas.Series` type object

```
wage = pd.Series([3.10, 3.24, 3.00, 6.00, 5.30, 8.75])  
print(wage)
```



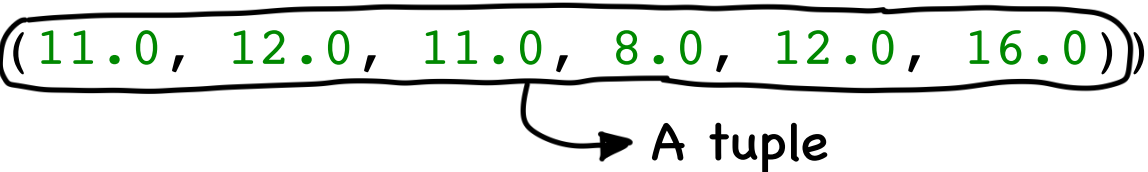
A list

```
0    3.10  
1    3.24  
2    3.00  
3    6.00  
4    5.30  
5    8.75  
dtype: float64
```

Pandas for Data Analysis

- The `pandas.Series` data structure
 - A one-dimensional array of indexed data
 - Create a `pandas.Series` type object

```
educ = pd.Series((11.0, 12.0, 11.0, 8.0, 12.0, 16.0))  
print(educ)
```




A tuple

```
0    11.0  
1    12.0  
2    11.0  
3     8.0  
4    12.0  
5    16.0  
dtype: float64
```

Pandas for Data Analysis

- The `pandas.Series` data structure
 - Attributes of series

```
wage = pd.Series([3.10, 3.24, 3.00, 6.00, 5.30, 8.75])  
print(wage)
```



Attributes of Charmander



HP:	39	<div></div>
Attack:	52	<div></div>
Defense:	43	<div></div>
Sp. Atk:	50	<div></div>
Sp. Def:	50	<div></div>
Speed:	65	<div></div>

Pandas for Data Analysis

- The `pandas.Series` data structure
 - Attributes of series

```
wage.values
```

```
array([3.1 , 3.24, 3.   , 6.   , 5.3 , 8.75])
```

Values of the series

wage

0	3.10
1	3.24
2	3.00
3	6.00
4	5.30
5	8.75

dtype: float64

Pandas for Data Analysis

- The `pandas.Series` data structure
 - Attributes of series

`wage.index`

`RangeIndex(start=0, stop=6, step=1)`

Indexes of the series

wage

0	3.10
1	3.24
2	3.00
3	6.00
4	5.30
5	8.75

`dtype: float64`

Pandas for Data Analysis

- The `pandas.Series` data structure

- ▶ Attributes of series

```
index = ['Mary', 'Ann', 'John', 'David', 'Frank', 'Ben']  
exper = pd.Series([2.0, 22.0, 2.0, 44.0, 7.0, 9.0],  
                  index=index)  
print(exper)
```

Specify indexes via the keyword argument

Mary	2.0
Ann	22.0
John	2.0
David	44.0
Frank	7.0
Ben	9.0
dtype: float64	

exper

Mary	2.0
Ann	22.0
John	2.0
David	44.0
Frank	7.0
Ben	9.0
dtype: float64	

Pandas for Data Analysis

- The `pandas.Series` data structure

- Attributes of series

```
print(exper.index)
```

```
Index(['Mary', 'Ann', 'John', 'David', 'Frank', 'Ben'], dtype='object')
```

Indexes of the series

exper

Mary	2.0
Ann	22.0
John	2.0
David	44.0
Frank	7.0
Ben	9.0
dtype: float64	

Pandas for Data Analysis

- The `pandas.Series` data structure
 - Attributes of series

```
print(educ.dtype)
```

float64

Type of data in the series

Pandas dtype	Built-in Python types
object	str or mixed types
int64	int
float64	float
bool	bool

educ

```
0    11.0
1    12.0
2    11.0
3     8.0
4    12.0
5    16.0
dtype: float64
```

Pandas for Data Analysis

- The `pandas.Series` data structure

- Attributes of series

```
educ_int = educ.astype(int)
print(educ_int)
```


Convert all data items to integers

```
0    11
1    12
2    11
3     8
4    12
5    16
dtype: int64
```

educ

0	11.0
1	12.0
2	11.0
3	8.0
4	12.0
5	16.0

dtype: float64



educ_int

0	11
1	12
2	11
3	8
4	12
5	16

dtype: int64

Pandas for Data Analysis




- The `pandas.Series` data structure

- Attributes of series

```
educ_str = educ.astype(str)  
print(educ_str)
```

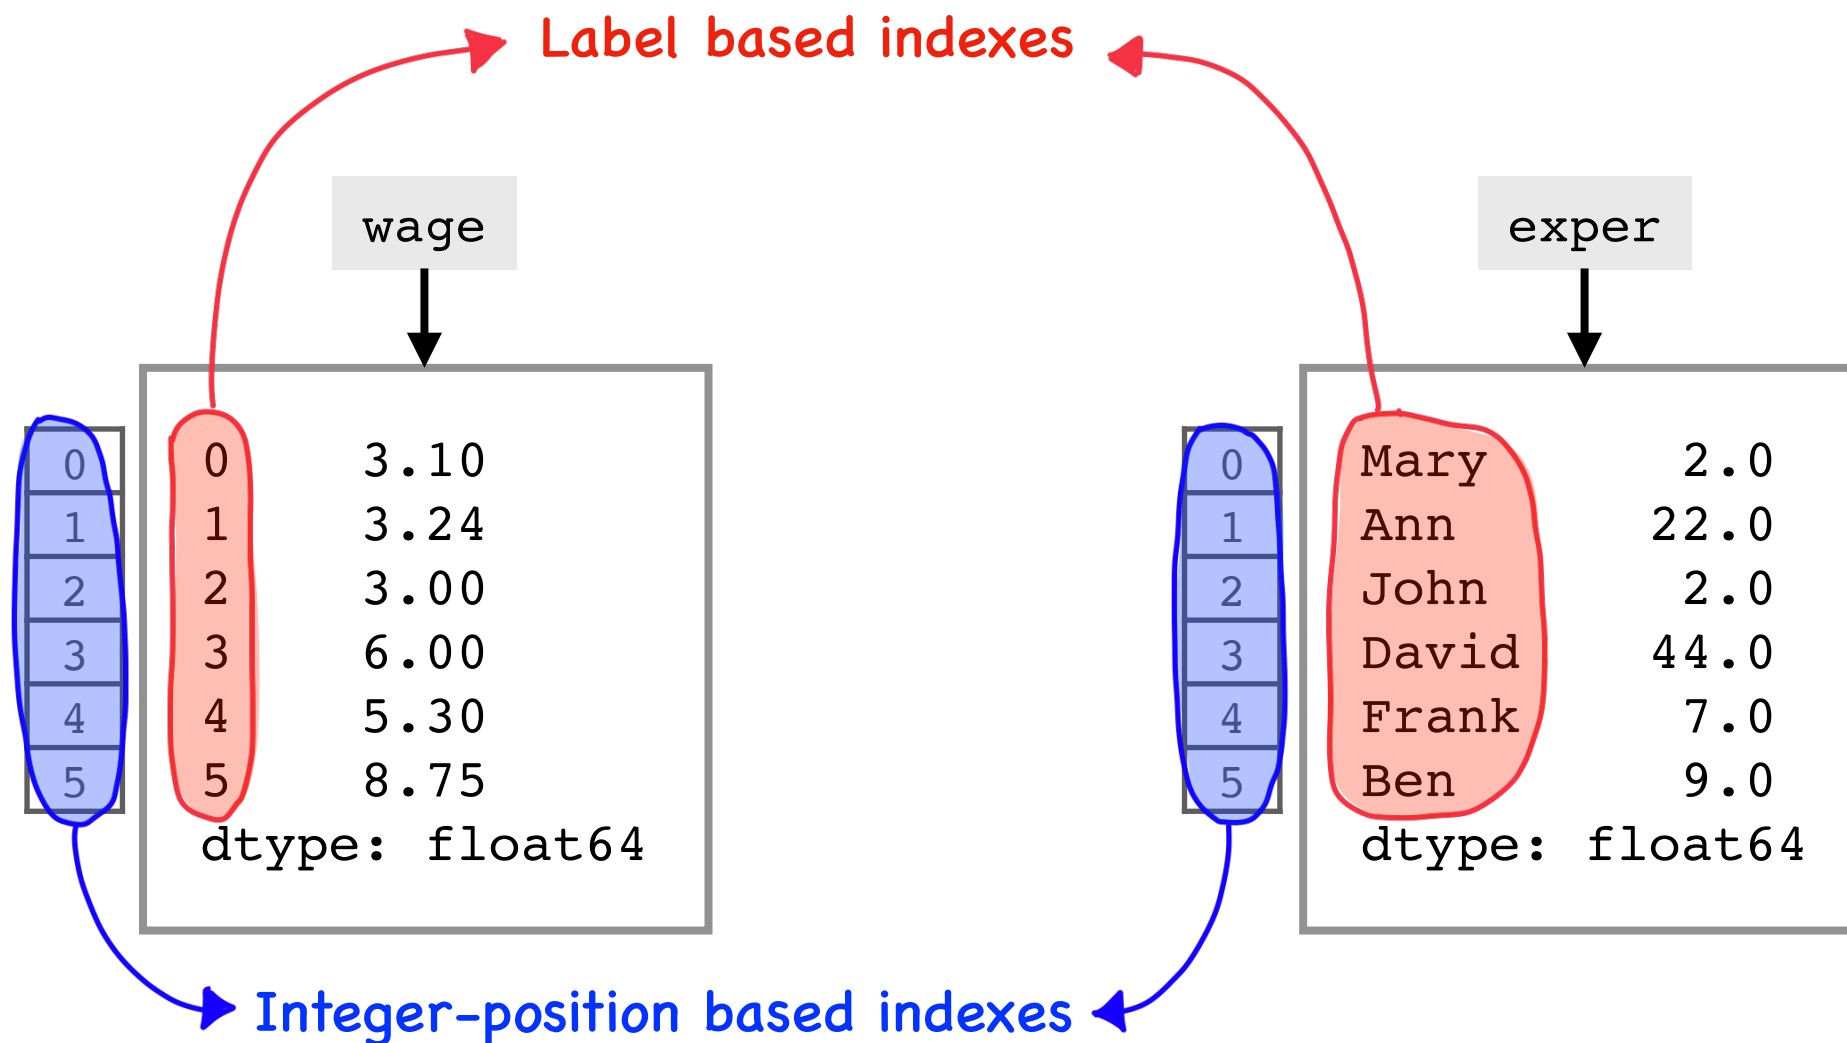
Convert all data items to strings

```
0    11.0  
1    12.0  
2    11.0  
3     8.0  
4    12.0  
5    16.0  
dtype: object
```



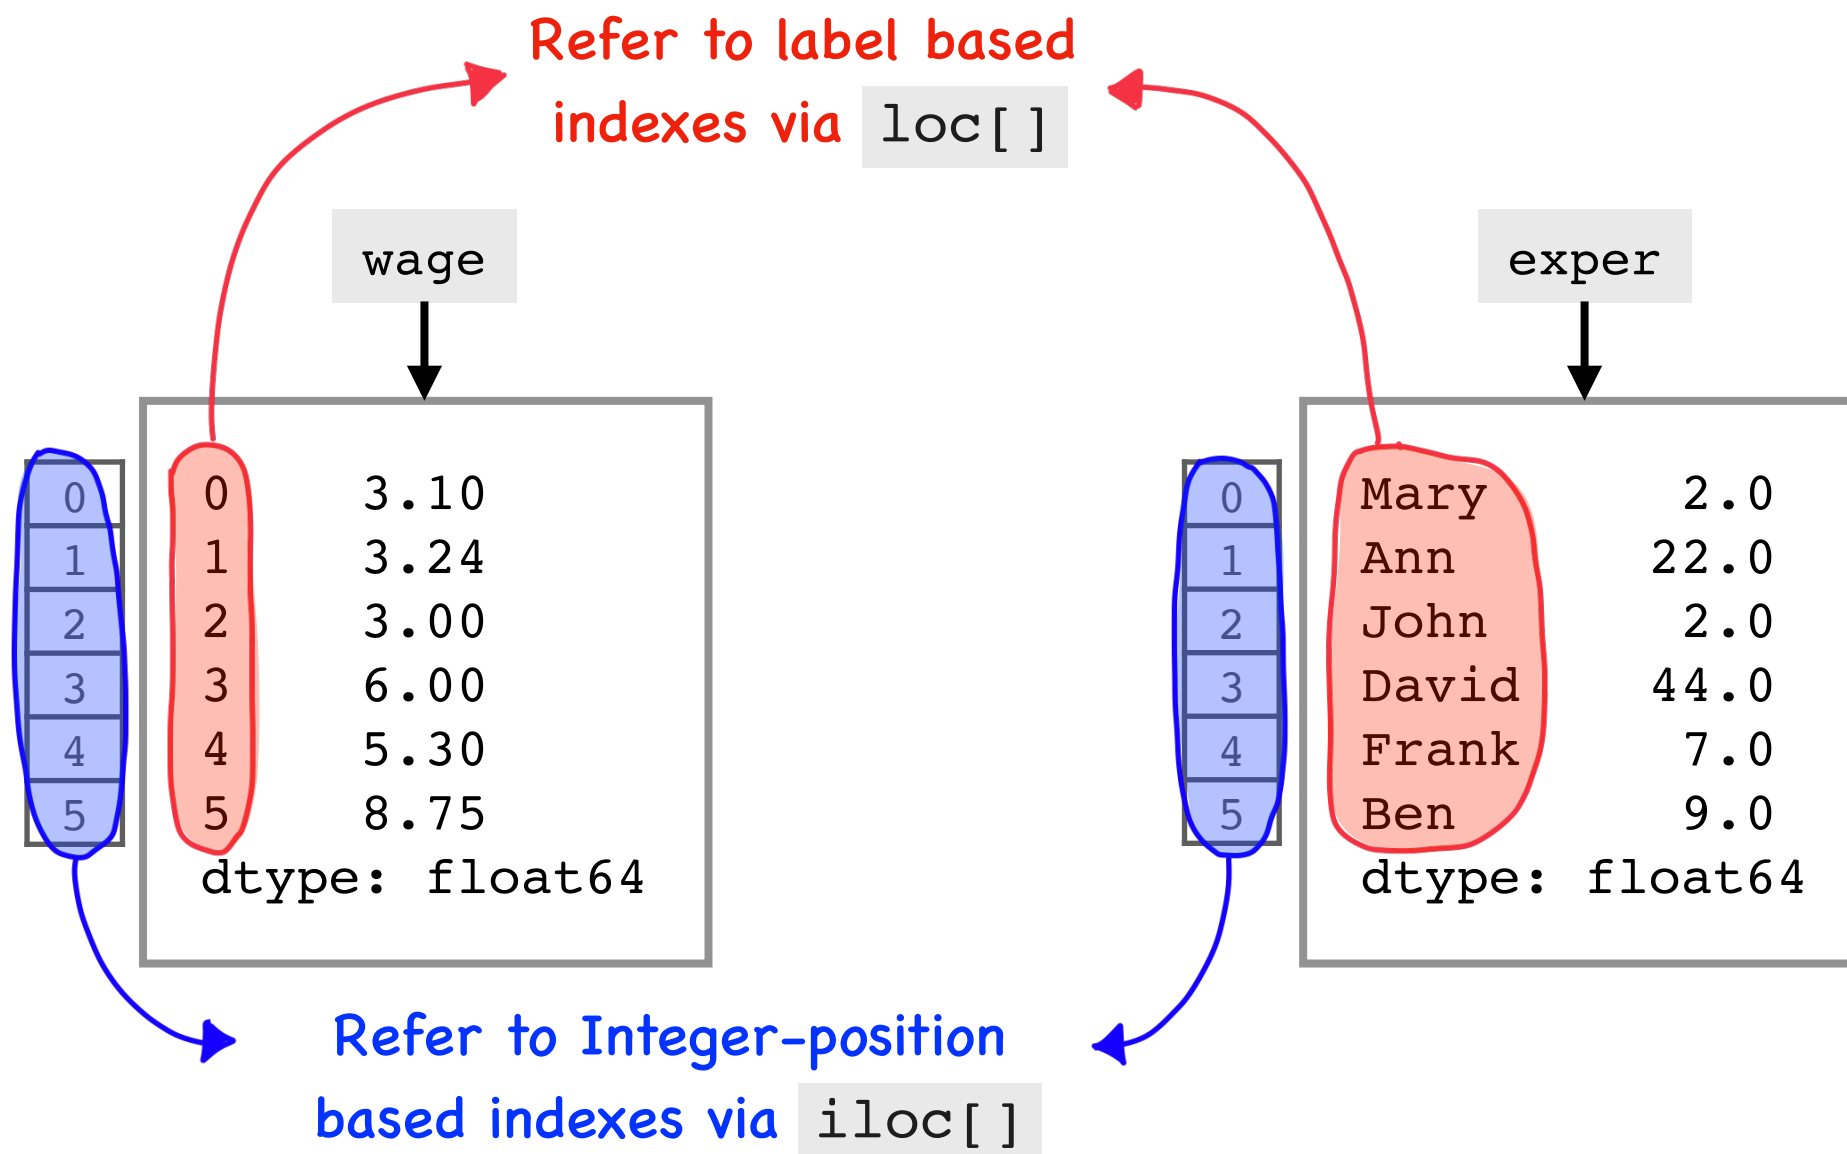
Pandas for Data Analysis

- The `pandas.Series` data structure
 - Indexing and slicing via `iloc[]` and `loc[]` indexers



Pandas for Data Analysis

- The `pandas.Series` data structure
 - Indexing and slicing via `iloc[]` and `loc[]` indexers



Pandas for Data Analysis

- The `pandas.Series` data structure
 - Indexing and slicing via `iloc[]` and `loc[]` indexers

```
print(exper.iloc[1])
```

```
print(exper.loc['Ann'])
```

22.0

exper

0	Mary	2.0
1	Ann	22.0
2	John	2.0
3	David	44.0
4	Frank	7.0
5	Ben	9.0

dtype: float64

Pandas for Data Analysis

- The `pandas.Series` data structure
 - Indexing and slicing via `iloc[]` and `loc[]` indexers

```
print(exper.iloc[:3])
```

```
print(exper.loc[:'John'])
```

```
Mary      2.0
Ann      22.0
John       2.0
dtype: float64
```

Stop index is excluded

Stop index is included

0
1
2
3
4
5

exper

Mary	2.0
Ann	22.0
John	2.0
David	44.0
Frank	7.0
Ben	9.0

dtype: float64

Notes: In the slicing expressions for label based indexes, the item indexed by **stop** are included in the selection.

Pandas for Data Analysis

- The `pandas.Series` data structure
 - Indexing and slicing via `iloc[]` and `loc[]` indexers

```
print(exper.iloc[[0, 2]])
```

```
print(exper.loc[['Mary', 'John']])
```

```
Mary    2.0  
John    2.0  
dtype: float64
```

exper

0	Mary	2.0
1	Ann	22.0
2	John	2.0
3	David	44.0
4	Frank	7.0
5	Ben	9.0

dtype: float64

Pandas for Data Analysis

- The `pandas.Series` data structure
 - Indexing and slicing via `iloc[]` and `loc[]` indexers

```
print(wage.iloc[:2])
```

```
0    3.10  
1    3.24  
Name: wage, dtype: float64
```

```
print(wage.loc[:2])
```

```
0    3.10  
1    3.24  
2    3.00  
Name: wage, dtype: float64
```

Stop index is excluded

Stop index is included

wage

0	3.10
1	3.24
2	3.00
3	6.00
4	5.30
5	8.75

dtype: float64

Pandas for Data Analysis

- The `pandas.DataFrame` data structure
 - Create a `pandas.DataFrame` object

```
data_dict = {'wage': [3.10, 3.24, 3.00, 6.00, 5.30, 8.75],  
            'educ': [11.0, 12.0, 11.0, 8.0, 12.0, 16.0],  
            'exper': [2.0, 22.0, 2.0, 44.0, 7.0, 9.0],  
            'gender': ['F', 'F', 'M', 'M', 'M', 'M'],  
            'married': [False, True, False, True, True, True]}
```

```
data_frame = pd.DataFrame(data_dict)
```

Function that creates a data frame

	wage	educ	exper	gender	married
0	3.10	11.0	2.0	F	False
1	3.24	12.0	22.0	F	True
2	3.00	11.0	2.0	M	False
3	6.00	8.0	44.0	M	True
4	5.30	12.0	7.0	M	True
5	8.75	16.0	9.0	M	True

Pandas for Data Analysis

- The `pandas.DataFrame` data structure
 - Create a `pandas.DataFrame` object

```
data_dict = {'wage': [3.10, 3.24, 3.00, 6.00, 5.30, 8.75],  
            'educ': [11.0, 12.0, 11.0, 8.0, 12.0, 16.0],  
            'exper': [2.0, 22.0, 2.0, 44.0, 7.0, 9.0],  
            'gender': ['F', 'F', 'M', 'M', 'M', 'M'],  
            'married': [False, True, False, True, True, True]}
```

Dictionary keys
become column labels

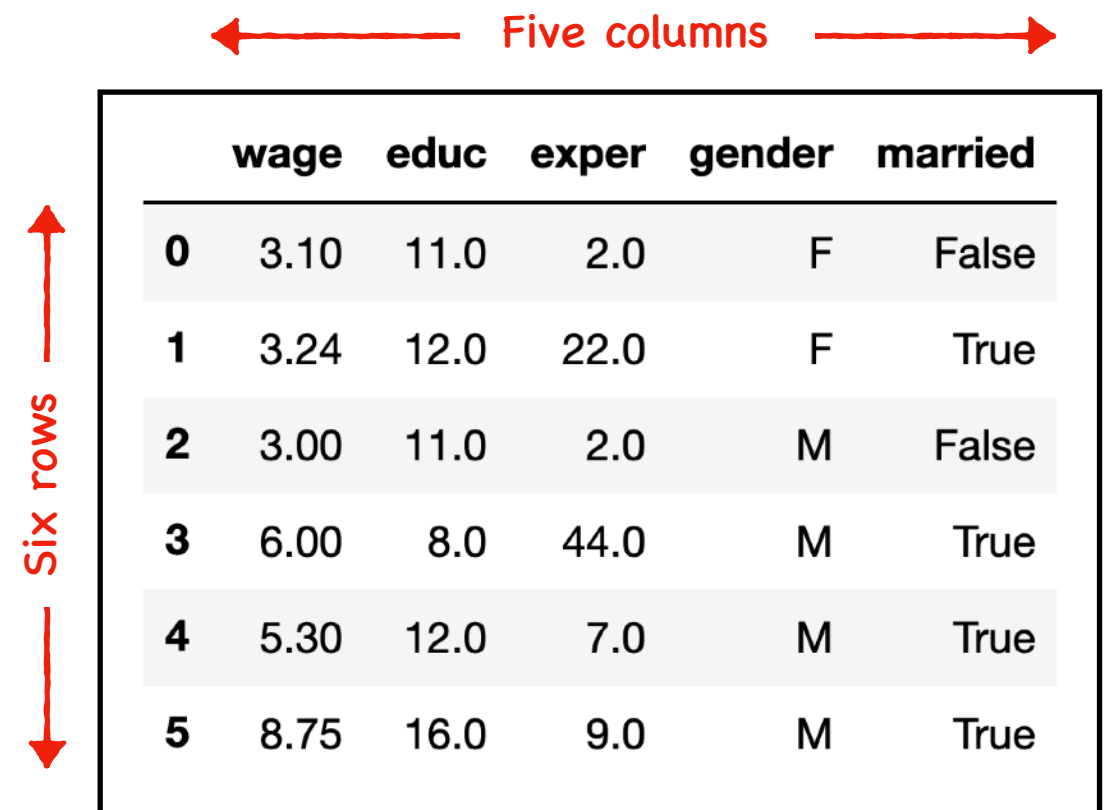
```
data_frame = pd.DataFrame(data_dict)
```

	wage	educ	exper	gender	married
0	3.10	11.0	2.0	F	False
1	3.24	12.0	22.0	F	True
2	3.00	11.0	2.0	M	False
3	6.00	8.0	44.0	M	True
4	5.30	12.0	7.0	M	True
5	8.75	16.0	9.0	M	True

Pandas for Data Analysis

- The `pandas.DataFrame` data structure
 - Attributes of data frames

Two-dimensional array



A diagram illustrating a pandas DataFrame as a two-dimensional array. It features a table with 6 rows and 5 columns. A red double-headed arrow above the table is labeled "Five columns". A red double-headed arrow to the left of the table is labeled "Six rows".

	wage	educ	exper	gender	married
0	3.10	11.0	2.0	F	False
1	3.24	12.0	22.0	F	True
2	3.00	11.0	2.0	M	False
3	6.00	8.0	44.0	M	True
4	5.30	12.0	7.0	M	True
5	8.75	16.0	9.0	M	True

Pandas for Data Analysis

- The `pandas.DataFrame` data structure

- ▶ Attributes of data frames

```
→ print(data_frame.columns)  
   print(data_frame.index)
```

```
Index(['wage', 'educ', 'exper', 'gender', 'married'], dtype='object')  
RangeIndex(start=0, stop=6, step=1)
```

	wage	educ	exper	gender	married
0	3.10	11.0	2.0	F	False
1	3.24	12.0	22.0	F	True
2	3.00	11.0	2.0	M	False
3	6.00	8.0	44.0	M	True
4	5.30	12.0	7.0	M	True
5	8.75	16.0	9.0	M	True

Pandas for Data Analysis

- The `pandas.DataFrame` data structure
 - Attributes of data frames

```
print(data_frame.columns)  
→ print(data_frame.index)
```

```
Index(['wage', 'educ', 'exper', 'gender', 'married'], dtype='object')  
RangeIndex(start=0, stop=6, step=1)
```

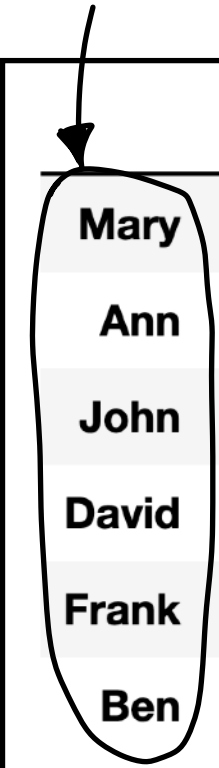
	wage	educ	exper	gender	married
0	3.10	11.0	2.0	F	False
1	3.24	12.0	22.0	F	True
2	3.00	11.0	2.0	M	False
3	6.00	8.0	44.0	M	True
4	5.30	12.0	7.0	M	True
5	8.75	16.0	9.0	M	True

Pandas for Data Analysis

- The `pandas.DataFrame` data structure
 - Attributes of data frames

```
index = ['Mary', 'Ann', 'John', 'David', 'Frank', 'Ben']  
data_frame_new = pd.DataFrame(data_dict, index=index)
```

Specify data frame row indexes via the keyword argument



	wage	educ	exper	gender	married
Mary	3.10	11.0	2.0	F	False
Ann	3.24	12.0	22.0	F	True
John	3.00	11.0	2.0	M	False
David	6.00	8.0	44.0	M	True
Frank	5.30	12.0	7.0	M	True
Ben	8.75	16.0	9.0	M	True

Pandas for Data Analysis

- The `pandas.DataFrame` data structure

- ▶ Attributes of data frames

```
→ print(data_frame_new.columns)  
print(data_frame_new.index)
```

```
Index(['wage', 'educ', 'exper', 'gender', 'married'], dtype='object')  
Index(['Mary', 'Ann', 'John', 'David', 'Frank', 'Ben'], dtype='object')
```

	wage	educ	exper	gender	married
Mary	3.10	11.0	2.0	F	False
Ann	3.24	12.0	22.0	F	True
John	3.00	11.0	2.0	M	False
David	6.00	8.0	44.0	M	True
Frank	5.30	12.0	7.0	M	True
Ben	8.75	16.0	9.0	M	True

Pandas for Data Analysis

- The `pandas.DataFrame` data structure
 - Attributes of data frames

```
print(data_frame_new.columns)  
→ print(data_frame_new.index)
```

```
Index(['wage', 'educ', 'exper', 'gender', 'married'], dtype='object')  
Index(['Mary', 'Ann', 'John', 'David', 'Frank', 'Ben'], dtype='object')
```

	wage	educ	exper	gender	married
Mary	3.10	11.0	2.0	F	False
Ann	3.24	12.0	22.0	F	True
John	3.00	11.0	2.0	M	False
David	6.00	8.0	44.0	M	True
Frank	5.30	12.0	7.0	M	True
Ben	8.75	16.0	9.0	M	True

Pandas for Data Analysis

- The `pandas.DataFrame` data structure
 - Attributes of data frames

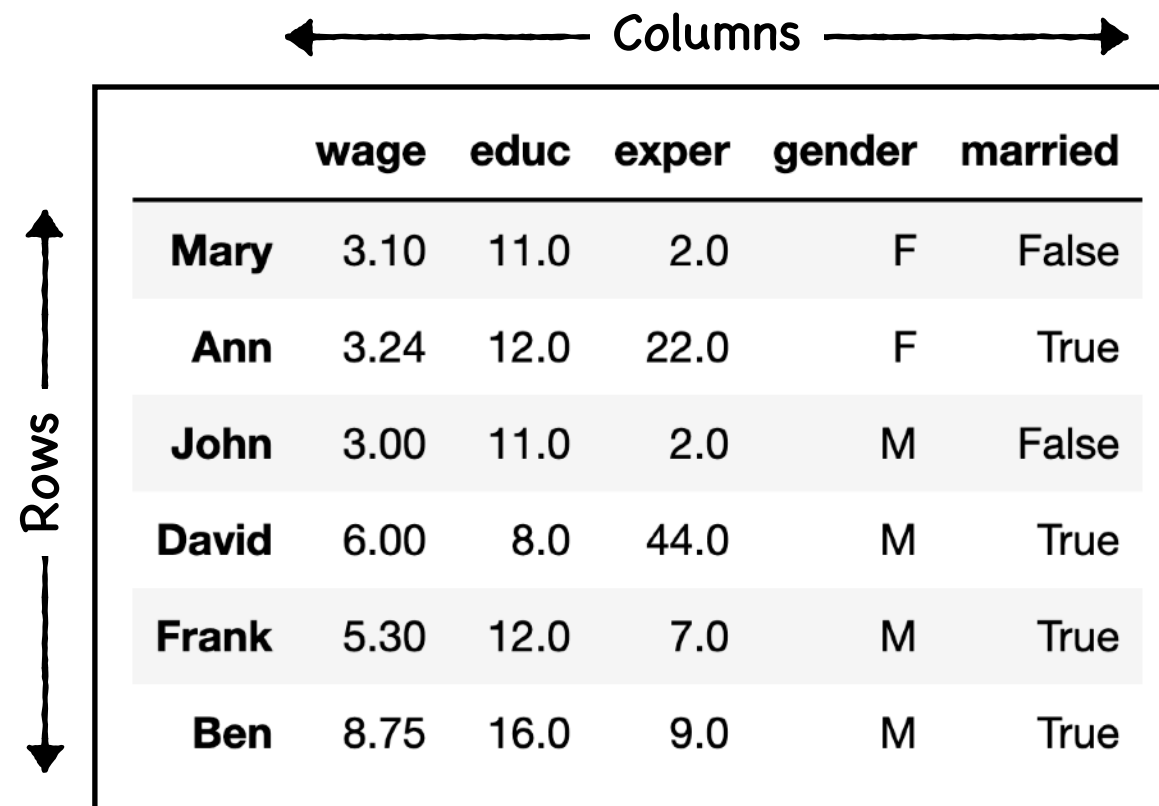
```
data_frame.dtypes
```

```
wage          float64
educ          float64
exper         float64
gender        object
married       bool
dtype: object
```

	wage	educ	exper	gender	married
Mary	3.10	11.0	2.0	F	False
Ann	3.24	12.0	22.0	F	True
John	3.00	11.0	2.0	M	False
David	6.00	8.0	44.0	M	True
Frank	5.30	12.0	7.0	M	True
Ben	8.75	16.0	9.0	M	True

Pandas for Data Analysis

- The `pandas.DataFrame` data structure
 - Indexing and slicing via `iloc[]` and `loc[]` indexers



A diagram illustrating a pandas DataFrame. It consists of a table with 6 rows and 6 columns. The columns are labeled 'wage', 'educ', 'exper', 'gender', and 'married'. The rows are labeled with names: 'Mary', 'Ann', 'John', 'David', 'Frank', and 'Ben'. A horizontal double-headed arrow above the table is labeled 'Columns'. A vertical double-headed arrow to the left of the table is labeled 'Rows'.

	wage	educ	exper	gender	married
Mary	3.10	11.0	2.0	F	False
Ann	3.24	12.0	22.0	F	True
John	3.00	11.0	2.0	M	False
David	6.00	8.0	44.0	M	True
Frank	5.30	12.0	7.0	M	True
Ben	8.75	16.0	9.0	M	True

Pandas for Data Analysis

- The `pandas.DataFrame` data structure
 - Indexing and slicing via `iloc[]` and `loc[]` indexers

Integer-position based indexes

	wage	educ	exper	gender	married	
0	Mary	3.10	11.0	2.0	F	False
1	Ann	3.24	12.0	22.0	F	True
2	John	3.00	11.0	2.0	M	False
3	David	6.00	8.0	44.0	M	True
4	Frank	5.30	12.0	7.0	M	True
5	Ben	8.75	16.0	9.0	M	True

Label based indexes

Pandas for Data Analysis

- The `pandas.DataFrame` data structure
 - Indexing and slicing via `iloc[]` and `loc[]` indexers

Selection using `iloc[rows, columns]`

0

1

2

3

4

0

1

2

3

4

5

0

1

2

3

4

	wage	educ	exper	gender	married
Mary	3.10	11.0	2.0	F	False
Ann	3.24	12.0	22.0	F	True
John	3.00	11.0	2.0	M	False
David	6.00	8.0	44.0	M	True
Frank	5.30	12.0	7.0	M	True
Ben	8.75	16.0	9.0	M	True

Row selections

- A single row index
- A slice of row indexes
- A list of row indexes

Row selections

- A single row index
- A slice of row indexes
- A list of row indexes

Column selections

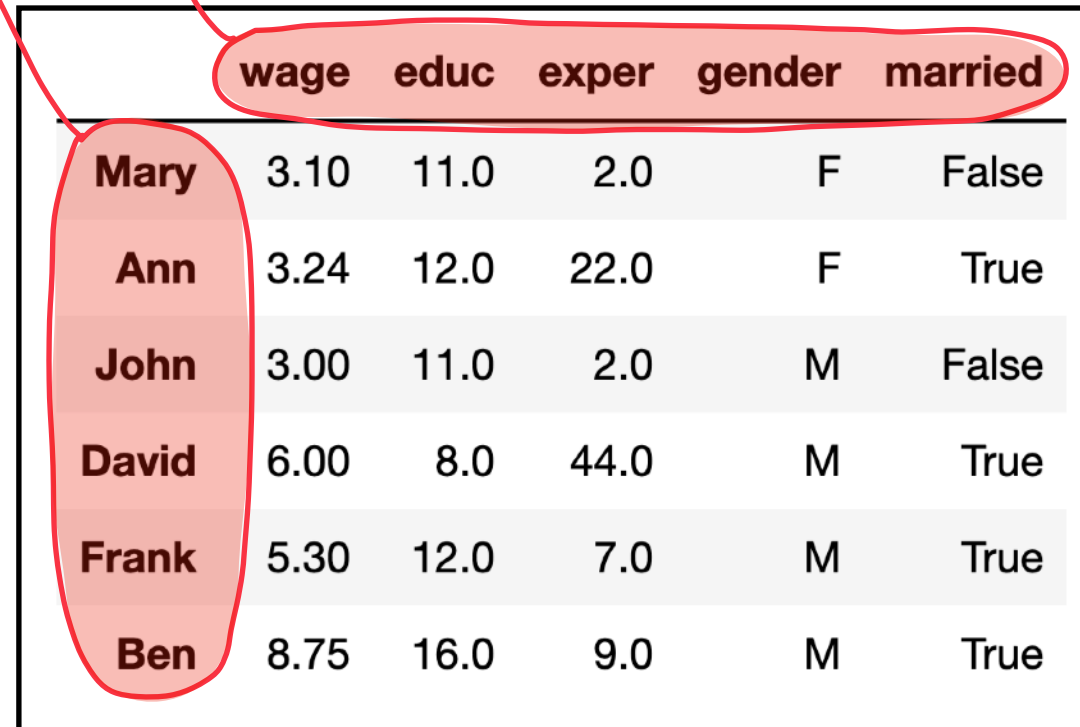
- A single column index
- A slice of column indexes
- A list of column indexes

The stop index is excluded from the row/column selection

Pandas for Data Analysis

- The `pandas.DataFrame` data structure
 - Indexing and slicing via `iloc[]` and `loc[]` indexers

Selection using `loc[rows, columns]`



	wage	educ	exper	gender	married
Mary	3.10	11.0	2.0	F	False
Ann	3.24	12.0	22.0	F	True
John	3.00	11.0	2.0	M	False
David	6.00	8.0	44.0	M	True
Frank	5.30	12.0	7.0	M	True
Ben	8.75	16.0	9.0	M	True

Row selections

- A single row label
- A slice of row labels
- A list of row labels
- Boolean type indexing

Column selections

- A single column label
- A slice of column labels
- A list of column labels

The stop index is included in the row/column selection

Pandas for Data Analysis

- The `pandas.DataFrame` data structure
 - Indexing and slicing via `iloc[]` and `loc[]` indexers

```
data_subset = data_frame.iloc[1:3, 1:3]  
print(data_subset)
```

Integer-position based column indexes

Integer-position based row indexes

Stop index is excluded

	educ	exper
1	12.0	22.0
2	11.0	2.0

Stop index is excluded

0
1
2
3
4
5

	0	1	2	3	4
wage	3.10	3.24	3.00	6.00	5.30
educ	11.0	12.0	11.0	8.0	12.0
exper	2.0	22.0	2.0	44.0	7.0
gender	F	F	M	M	M
married	False	True	False	True	True

Pandas for Data Analysis

- The `pandas.DataFrame` data structure
 - Indexing and slicing via `iloc[]` and `loc[]` indexers

```
data_subset = data_frame_new.iloc[1:3, 1:3]  
print(data_subset)
```

Integer-position based column indexes

	educ	exper
Ann	12.0	22.0
John	11.0	2.0

Integer-position based row indexes

Stop index is excluded

Stop index is excluded

	0	1	2	3	4
	wage	educ	exper	gender	married
Mary	3.10	11.0	2.0	F	False
Ann	3.24	12.0	22.0	F	True
John	3.00	11.0	2.0	M	False
David	6.00	8.0	44.0	M	True
Frank	5.30	12.0	7.0	M	True
Ben	8.75	16.0	9.0	M	True

Pandas for Data Analysis

- The `pandas.DataFrame` data structure
 - Indexing and slicing via `iloc[]` and `loc[]` indexers

```
cols = data_frame.iloc[:, 2]  
print(cols)
```

Integer-position based column index

All rows

```
0      2.0  
1     22.0  
2      1.0  
3     44.0  
4      7.0  
5      9.0  
Name: exper, dtype: float64
```

		0	1	2	3	4
		wage	educ	exper	gender	married
0	Mary	3.10	11.0	2.0	F	False
1	Ann	3.24	12.0	22.0	F	True
2	John	3.00	11.0	2.0	M	False
3	David	6.00	8.0	44.0	M	True
4	Frank	5.30	12.0	7.0	M	True
5	Ben	8.75	16.0	9.0	M	True

Pandas for Data Analysis

- The `pandas.DataFrame` data structure
 - Indexing and slicing via `iloc[]` and `loc[]` indexers

```
rows = data_frame_new.iloc[2:5, :]  
print(rows)
```

All columns

	wage	educ	exper	gender	married
John	3.0	11.0	2.0	M	False
David	6.0	8.0	44.0	M	True
Frank	5.3	12.0	7.0	M	True

Integer-position based row indexes

		0	1	2	3	4
		wage	educ	exper	gender	married
0	Mary	3.10	11.0	2.0	F	False
1	Ann	3.24	12.0	22.0	F	True
2	John	3.00	11.0	2.0	M	False
3	David	6.00	8.0	44.0	M	True
4	Frank	5.30	12.0	7.0	M	True
5	Ben	8.75	16.0	9.0	M	True

Pandas for Data Analysis

- The `pandas.DataFrame` data structure
 - Indexing and slicing via `iloc[]` and `loc[]` indexers

```
rows = data_frame_new.iloc[2:5]  
print(rows)
```

	wage	educ	exper	gender	married
John	3.0	11.0	2.0	M	False
David	6.0	8.0	44.0	M	True
Frank	5.3	12.0	7.0	M	True

Integer-position based row indexes

0
1
2
3
4
5

All columns by default

	0	1	2	3	4
	wage	educ	exper	gender	married
Mary	3.10	11.0	2.0	F	False
Ann	3.24	12.0	22.0	F	True
John	3.00	11.0	2.0	M	False
David	6.00	8.0	44.0	M	True
Frank	5.30	12.0	7.0	M	True
Ben	8.75	16.0	9.0	M	True

Pandas for Data Analysis

- The `pandas.DataFrame` data structure
 - Indexing and slicing via `iloc[]` and `loc[]` indexers

```
data_subset = data_frame.loc[1:2, 'educ': 'exper']  
print(data_subset)
```

	educ	exper
1	12.0	22.0
2	11.0	2.0

Label based row indexes

Label based column indexes

Stop index is included

Stop index is included

	wage	educ	exper	gender	married
0	3.10	11.0	2.0	F	False
1	3.24	12.0	22.0	F	True
2	3.00	11.0	2.0	M	False
3	6.00	8.0	44.0	M	True
4	5.30	12.0	7.0	M	True
5	8.75	16.0	9.0	M	True

Pandas for Data Analysis

- The `pandas.DataFrame` data structure
 - Indexing and slicing via `iloc[]` and `loc[]` indexers

```
data_new_subset = data_frame_new.loc['Ann':'John', 'educ':'exper']  
print(data_new_subset)
```

	educ	exper
Ann	12.0	22.0
John	11.0	2.0

Label based row indexes

Label based column indexes

Stop index is included

Stop index is included

	wage	educ	exper	gender	married
Mary	3.10	11.0	2.0	F	False
Ann	3.24	12.0	22.0	F	True
John	3.00	11.0	2.0	M	False
David	6.00	8.0	44.0	M	True
Frank	5.30	12.0	7.0	M	True
Ben	8.75	16.0	9.0	M	True

Pandas for Data Analysis

- The `pandas.DataFrame` data structure
 - Indexing and slicing via `iloc[]` and `loc[]` indexers

```
data_new_subset = data_frame_new.loc['Ann':'John', 'educ':'exper']  
print(data_new_subset)
```

	educ	exper
Ann	12.0	22.0
John	11.0	2.0

Label based row indexes

Label based column indexes

Notes: In the slicing expressions for label based indexes, the item indexed by **stop** are included in the selection.

	wage	educ	exper	gender	married
Mary	3.10	11.0	2.0	F	False
Ann	3.24	12.0	22.0	F	True
John	3.00	11.0	2.0	M	False
David	6.00	8.0	44.0	M	True
Frank	5.30	12.0	7.0	M	True
Ben	8.75	16.0	9.0	M	True

Pandas for Data Analysis

- The `pandas.DataFrame` data structure
 - Indexing and slicing via `iloc[]` and `loc[]` indexers

```
cols = data_frame.loc[:, 'exper']  
print(cols)
```

```
0      2.0  
1     22.0  
2      1.0  
3     44.0  
4      7.0  
5      9.0  
Name: exper, dtype: float64
```

All rows

Label based column indexes

	wage	educ	exper	gender	married
0	3.10	11.0	2.0	F	False
1	3.24	12.0	22.0	F	True
2	3.00	11.0	2.0	M	False
3	6.00	8.0	44.0	M	True
4	5.30	12.0	7.0	M	True
5	8.75	16.0	9.0	M	True

Pandas for Data Analysis

- The `pandas.DataFrame` data structure
 - Indexing and slicing via `iloc[]` and `loc[]` indexers

```
rows = data_frame_new.loc['John':'Frank', :]  
print(rows)
```

	wage	educ	exper	gender	married
John	3.00	11.0	2.0	M	False
David	6.00	8.0	44.0	M	True
Frank	5.30	12.0	7.0	M	True

Label based row indexes

All columns

	wage	educ	exper	gender	married
Mary	3.10	11.0	2.0	F	False
Ann	3.24	12.0	22.0	F	True
John	3.00	11.0	2.0	M	False
David	6.00	8.0	44.0	M	True
Frank	5.30	12.0	7.0	M	True
Ben	8.75	16.0	9.0	M	True

Pandas for Data Analysis

- The `pandas.DataFrame` data structure
 - Indexing and slicing via `iloc[]` and `loc[]` indexers

```
rows = data_frame_new.loc['John':'Frank']  
print(rows)
```

	wage	educ	exper	gender	married
John	3.00	11.0	2.0	M	False
David	6.00	8.0	44.0	M	True
Frank	5.30	12.0	7.0	M	True

Label based row indexes

All columns by default

	wage	educ	exper	gender	married
Mary	3.10	11.0	2.0	F	False
Ann	3.24	12.0	22.0	F	True
John	3.00	11.0	2.0	M	False
David	6.00	8.0	44.0	M	True
Frank	5.30	12.0	7.0	M	True
Ben	8.75	16.0	9.0	M	True

Pandas for Data Analysis

- The `pandas.DataFrame` data structure
 - Indexing of columns

```
data_frame[ 'educ' ]
```

```
0      11.0
1      12.0
2      11.0
3       8.0
4      12.0
5      16.0
Name: educ, dtype: float64
```

```
data_frame_new[ 'educ' ]
```

```
Mary      11.0
Ann       12.0
John      11.0
David      8.0
Frank     12.0
Ben       16.0
Name: educ, dtype: float64
```

Pandas for Data Analysis

- The `pandas.DataFrame` data structure
 - Change `pandas.DataFrame` object in-place
 - ✓ Similar syntax as other mutable data types

Pandas for Data Analysis

- The `pandas.DataFrame` data structure
 - Change `pandas.DataFrame` object in-place

```
data_frame.loc[2:3, 'educ'] = 9.0
```

	wage	educ	exper	gender	married
0	3.10	11.0	2.0	F	False
1	3.24	12.0	22.0	F	True
2	3.00	11.0	2.0	M	False
3	6.00	8.0	44.0	M	True
4	5.30	12.0	7.0	M	True
5	8.75	16.0	9.0	M	True



	wage	educ	exper	gender	married
0	3.10	11.0	2.0	F	False
1	3.24	12.0	22.0	F	True
2	3.00	9.0	2.0	M	False
3	6.00	9.0	44.0	M	True
4	5.30	12.0	7.0	M	True
5	8.75	16.0	9.0	M	True

Pandas for Data Analysis

- The `pandas.DataFrame` data structure
 - Change `pandas.DataFrame` object in-place

```
data_frame.iloc[2, 1:3] = 1.0
```

0

1

2

3

4

0

1

2

3

4

5

wage

educ

exper

gender

married

0

3.10

11.0

2.0

F

False

1

3.24

12.0

22.0

F

True

2

3.00

11.0

2.0

M

False

3

6.00

8.0

44.0

M

True

4

5.30

12.0

7.0

M

True

5

8.75

16.0

9.0

M

True

0

1

2

3

4

wage

educ

exper

gender

married

0

3.10

11.0

2.0

F

False

1

3.24

12.0

22.0

F

True

2

3.00

1.0

1.0

M

False

3

6.00

9.0

44.0

M

True

4

5.30

12.0

7.0

M

True

5

8.75

16.0

9.0

M

True

0

1

2

3

4

wage

educ

exper

gender

married

0

3.10

11.0

2.0

F

False

1

3.24

12.0

22.0

F

True

2

3.00

1.0

1.0

M

False

3

6.00

9.0

44.0

M

True

4

5.30

12.0

7.0

M

True

5

8.75

16.0

9.0

M

True

Pandas for Data Analysis

- The `pandas.DataFrame` data structure
 - Change `pandas.DataFrame` object in-place

```
data_frame.loc[:, 'remarks'] = 'none'
```


	wage	educ	exper	gender	married
0	3.10	11.0	2.0	F	False
1	3.24	12.0	22.0	F	True
2	3.00	1.0	1.0	M	False
3	6.00	9.0	44.0	M	True
4	5.30	12.0	7.0	M	True
5	8.75	16.0	9.0	M	True

Pandas for Data Analysis

- The `pandas.DataFrame` data structure
 - Change `pandas.DataFrame` object in-place

```
data_frame.loc[:, 'remarks'] = 'none'
```

	wage	educ	exper	gender	married
0	3.10	11.0	2.0	F	False
1	3.24	12.0	22.0	F	True
2	3.00	1.0	1.0	M	False
3	6.00	9.0	44.0	M	True
4	5.30	12.0	7.0	M	True
5	8.75	16.0	9.0	M	True




	wage	educ	exper	gender	married	remarks
0	3.10	11.0	2.0	F	False	none
1	3.24	12.0	22.0	F	True	none
2	3.00	1.0	1.0	M	False	none
3	6.00	9.0	44.0	M	True	none
4	5.30	12.0	7.0	M	True	none
5	8.75	16.0	9.0	M	True	none

Pandas for Data Analysis

- The `pandas.DataFrame` data structure
 - Boolean series and boolean indexing

	wage	educ	exper	gender	married	remarks
0	3.10	11.0	2.0	F	False	none
1	3.24	12.0	22.0	F	True	none
2	3.00	1.0	1.0	M	False	none
3	6.00	9.0	44.0	M	True	none
4	5.30	12.0	7.0	M	True	none
5	8.75	16.0	9.0	M	True	none



Select all married cases

	wage	educ	exper	gender	married	remarks
1	3.24	12.0	22.0	F	True	none
3	6.00	9.0	44.0	M	True	none
4	5.30	12.0	7.0	M	True	none
5	8.75	16.0	9.0	M	True	none

Pandas for Data Analysis

- The `pandas.DataFrame` data structure
 - Boolean series and boolean indexing

```
is_female = data_frame['gender'] == 'F'  
print(is_female)
```

```
0      True  
1      True  
2     False  
3     False  
4     False  
5     False  
Name: gender, dtype: bool
```

`data_frame['gender']`

0	F	== 'F'
1	F	== 'F'
2	M	== 'F'
3	M	== 'F'
4	M	== 'F'
5	M	== 'F'

dtype: object



`is_female`

0	True
1	True
2	False
3	False
4	False
5	False

dtype: bool

Pandas for Data Analysis

- The `pandas.DataFrame` data structure
 - Boolean series and boolean indexing

```
is_high_wage = data_frame['wage'] > 4  
print(is_high_wage)
```

```
0    False  
1    False  
2    False  
3     True  
4     True  
5     True  
Name: wage, dtype: bool
```

data_frame['wage']

0	3.10	> 4
1	3.24	> 4
2	3.00	> 4
3	6.00	> 4
4	5.30	> 4
5	8.75	> 4

dtype: float64



is_high_wage

0	False
1	False
2	False
3	True
4	True
5	True

dtype: bool

Pandas for Data Analysis

- The `pandas.DataFrame` data structure
 - Boolean series and boolean indexing

```
not_female = data_frame['gender'] != 'F'  
print(not_female)
```

```
0    False  
1    False  
2     True  
3     True  
4     True  
5     True  
Name: gender, dtype: bool
```

`data_frame['gender']`

0	F	!= 'F'
1	F	!= 'F'
2	M	!= 'F'
3	M	!= 'F'
4	M	!= 'F'
5	M	!= 'F'

dtype: object



`not_female`

0	False
1	False
2	True
3	True
4	True
5	True

dtype: bool

Pandas for Data Analysis

- The `pandas.DataFrame` data structure
 - Boolean series and boolean indexing
 - ✓ Bitwise “and” logic: `&`
 - ✓ Bitwise “or” logic: `|`
 - ✓ Bitwise “not” logic: `~`

Pandas for Data Analysis

- The `pandas.DataFrame` data structure
 - Boolean series and boolean indexing

```
→ cond1 = data_frame['gender'] == 'F'  
cond2 = data_frame['married']  
is_wife = cond1 & cond2  
print(is_wife)
```

```
0    False  
1     True  
2    False  
3    False  
4    False  
5    False  
dtype: bool
```

data_frame['gender']

0	F	== 'F'
1	F	== 'F'
2	M	== 'F'
3	M	== 'F'
4	M	== 'F'
5	M	== 'F'

dtype: object

cond1

0	True
1	True
2	False
3	False
4	False
5	False

dtype: bool

Pandas for Data Analysis

- The `pandas.DataFrame` data structure
 - Boolean series and boolean indexing

```
cond1 = data_frame['gender'] == 'F'  
→ cond2 = data_frame['married']  
is_wife = cond1 & cond2  
print(is_wife)
```

```
0    False  
1     True  
2    False  
3    False  
4    False  
5    False  
dtype: bool
```

cond2

↓

```
0    False  
1     True  
2    False  
3     True  
4     True  
5     True  
dtype: bool
```

Pandas for Data Analysis

- The `pandas.DataFrame` data structure
 - Boolean series and boolean indexing

```
cond1 = data_frame['gender'] == 'F'  
cond2 = data_frame['married']  
→ is_wife = cond1 & cond2  
print(is_wife)
```

```
0    False  
1     True  
2    False  
3    False  
4    False  
5    False  
dtype: bool
```

cond1

```
0    True  
1    True  
2   False  
3   False  
4   False  
5   False  
dtype: bool
```

&

cond2

```
0   False  
1    True  
2   False  
3    True  
4    True  
5    True  
dtype: bool
```

Pandas for Data Analysis

- The `pandas.DataFrame` data structure
 - Boolean series and boolean indexing

```
cond1 = data_frame['gender'] == 'F'  
cond2 = data_frame['married']  
→ is_wife = cond1 & cond2  
print(is_wife)
```

```
0    False  
1     True  
2    False  
3    False  
4    False  
5    False  
dtype: bool
```

cond1

```
0    True  
1    True  
2    False  
3    False  
4    False  
5    False  
dtype: bool
```

and
and
and
and
and

cond2

```
0    False  
1     True  
2    False  
3     True  
4     True  
5     True  
dtype: bool
```



is_wife

```
0    False  
1     True  
2    False  
3    False  
4    False  
5    False  
dtype: bool
```


Pandas for Data Analysis

- The `pandas.DataFrame` data structure
 - Boolean series and boolean indexing

```
→ cond1 = data_frame['educ'] > 9
   cond2 = data_frame['exper'] > 3
   is_skillful = cond1 | cond2
   print(is_skillful)
```

```
0      True
1      True
2     False
3      True
4      True
5      True
dtype: bool
```

cond1

```
0      True
1      True
2     False
3     False
4      True
5      True
dtype: bool
```

cond2

```
0     False
1      True
2     False
3      True
4      True
5      True
dtype: bool
```

Pandas for Data Analysis

- The `pandas.DataFrame` data structure
 - Boolean series and boolean indexing

```
cond1 = data_frame['educ'] > 9
cond2 = data_frame['exper'] > 3
→ is_skillful = cond1 | cond2
print(is_skillful)
```

0	True
1	True
2	False
3	True
4	True
5	True

dtype: bool

cond1



0	True
1	True
2	False
3	False
4	True
5	True

dtype: bool

or
or
or
or
or
or

cond2



0	False
1	True
2	False
3	True
4	True
5	True

dtype: bool



is_skillful



0	True
1	True
2	False
3	True
4	True
5	True

dtype: bool

Pandas for Data Analysis

- The `pandas.DataFrame` data structure
 - Boolean series and boolean indexing

```
→ is_female = data_frame['gender'] == 'F'  
females = data_frame.loc[is_female]
```

is_female

0	True
1	True
2	False
3	False
4	False
5	False

dtype: bool

data_frame

	wage	educ	exper	gender	married	remarks
0	3.10	11.0	2.0	F	False	none
1	3.24	12.0	22.0	F	True	none
2	3.00	1.0	1.0	M	False	none
3	6.00	9.0	44.0	M	True	none
4	5.30	12.0	7.0	M	True	none
5	8.75	16.0	9.0	M	True	none

Pandas for Data Analysis

- The `pandas.DataFrame` data structure
 - Boolean series and boolean indexing

```
is_female = data_frame['gender'] == 'F'  
→ females = data_frame.loc[is_female]
```

is_female

```
0    True  
1    True  
2   False  
3   False  
4   False  
5   False  
dtype: bool
```

females

	wage	educ	exper	gender	married	remarks
0	3.10	11.0	2.0	F	False	none
1	3.24	12.0	22.0	F	True	none

Pandas for Data Analysis

- The `pandas.DataFrame` data structure
 - Boolean series and boolean indexing

```
cond1 = data_frame['gender'] == 'F'  
cond2 = data_frame['married']  
→ is_wife = cond1 & cond2  
data_frame.loc[is_wife, 'remarks'] = 'Wife'
```

is_wife

0	False
1	True
2	False
3	False
4	False
5	False
dtype: bool	

data_frame

	wage	educ	exper	gender	married	remarks
0	3.10	11.0	2.0	F	False	none
1	3.24	12.0	22.0	F	True	none
2	3.00	1.0	1.0	M	False	none
3	6.00	9.0	44.0	M	True	none
4	5.30	12.0	7.0	M	True	none
5	8.75	16.0	9.0	M	True	none

Pandas for Data Analysis

- The `pandas.DataFrame` data structure
 - Boolean series and boolean indexing

```
cond1 = data_frame['gender'] == 'F'  
cond2 = data_frame['married']  
is_wife = cond1 & cond2  
→ data_frame.loc[is_wife, 'remarks'] = 'Wife'
```

is_wife

0	False
1	True
2	False
3	False
4	False
5	False

dtype: bool

data_frame

	wage	educ	exper	gender	married	remarks
0	3.10	11.0	2.0	F	False	none
1	3.24	12.0	22.0	F	True	Wife
2	3.00	1.0	1.0	M	False	none
3	6.00	9.0	44.0	M	True	none
4	5.30	12.0	7.0	M	True	none
5	8.75	16.0	9.0	M	True	none

Pandas for Data Analysis

- The `pandas.DataFrame` data structure
 - Boolean series and boolean indexing

Question 1: Make changes to the data frame above such that:

- Values for married males in the column '**remarks**' are changed to the string '**Husband**'.
- Values for unmarried males or unmarried females in the column '**remarks**' are changed to the string '**Single**'.

Pandas for Data Analysis

- The `pandas.DataFrame` data structure
 - Boolean series and boolean indexing

Question 1: Make changes to the data frame above such that:

- Values for married males in the column '**remarks**' are changed to the string '**Husband**'.
- Values for unmarried males or unmarried females in the column '**remarks**' are changed to the string '**Single**'.

```
cond1 = data_frame['gender'] == 'F'
cond2 = data_frame['married']
is_husband = ~cond1 & cond2
data_frame.loc[is_husband, 'remarks'] = 'Husband'
is_single = ~cond2
data_frame.loc[is_single, 'remarks'] = 'Single'
```


Pandas for Data Analysis

- Read data from files

```
data = pd.read_csv('wage.csv')
```

```
print(data.head(6))
```

The first six records

	wage	educ	exper	gender	married
0	3.10	11.0	2.0	F	False
1	3.24	12.0	22.0	F	True
2	3.00	11.0	2.0	M	False
3	6.00	8.0	44.0	M	True
4	5.30	12.0	7.0	M	True
5	8.75	16.0	9.0	M	True

	wage	educ	exper	gender	married
0	3.10	11.0	2.0	F	False
1	3.24	12.0	22.0	F	True
2	3.00	11.0	2.0	M	False
3	6.00	8.0	44.0	M	True
4	5.30	12.0	7.0	M	True
...
521	15.00	16.0	14.0	F	True
522	2.27	10.0	2.0	F	False
523	4.67	15.0	13.0	M	True
524	11.56	16.0	5.0	M	True
525	3.50	14.0	5.0	F	False

526 rows x 5 columns

Basics of Descriptive Analytics

- Descriptive measures

```
data_num = data.drop(columns='gender')
```

Drop the selected column

data

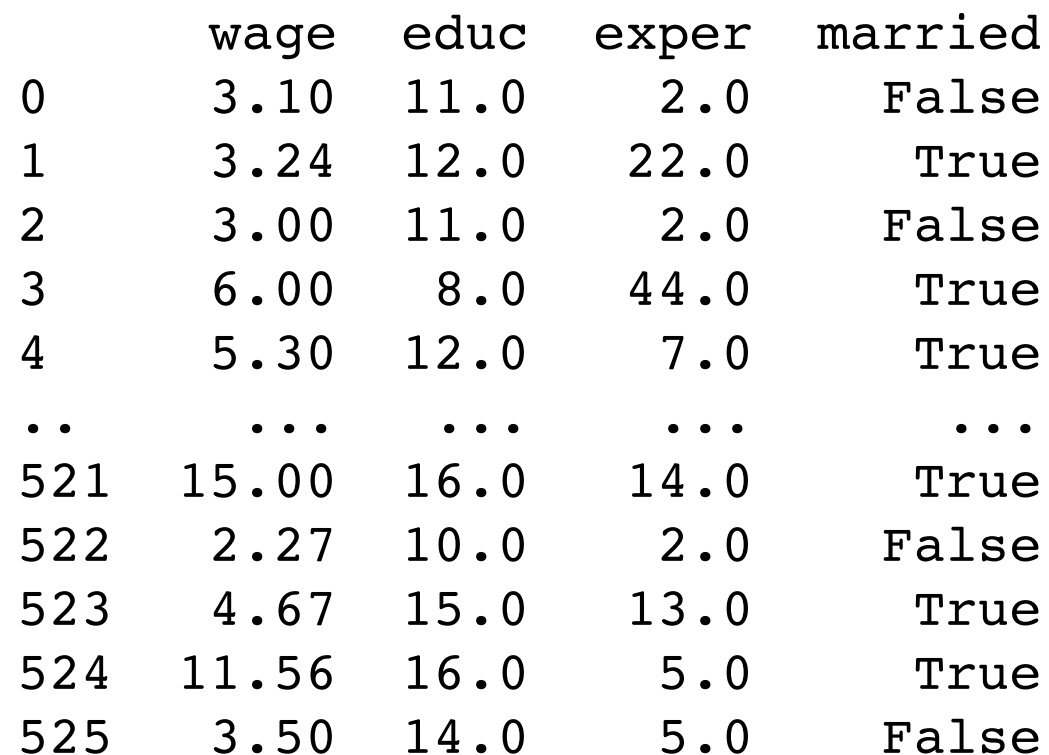
	wage	educ	exper	gender	married
0	3.10	11.0	2.0	F	False
1	3.24	12.0	22.0	F	True
2	3.00	11.0	2.0	M	False
3	6.00	8.0	44.0	M	True
4	5.30	12.0	7.0	M	True
..
521	15.00	16.0	14.0	F	True
522	2.27	10.0	2.0	F	False
523	4.67	15.0	13.0	M	True
524	11.56	16.0	5.0	M	True
525	3.50	14.0	5.0	F	False

Basics of Descriptive Analytics

- Descriptive measures

```
data_num = data.drop(columns='gender')
```

data_num



	wage	educ	exper	married
0	3.10	11.0	2.0	False
1	3.24	12.0	22.0	True
2	3.00	11.0	2.0	False
3	6.00	8.0	44.0	True
4	5.30	12.0	7.0	True
..
521	15.00	16.0	14.0	True
522	2.27	10.0	2.0	False
523	4.67	15.0	13.0	True
524	11.56	16.0	5.0	True
525	3.50	14.0	5.0	False

Basics of Descriptive Analytics


- Descriptive measures

- Measures of centers

- ✓ Mean

- ✓ Median

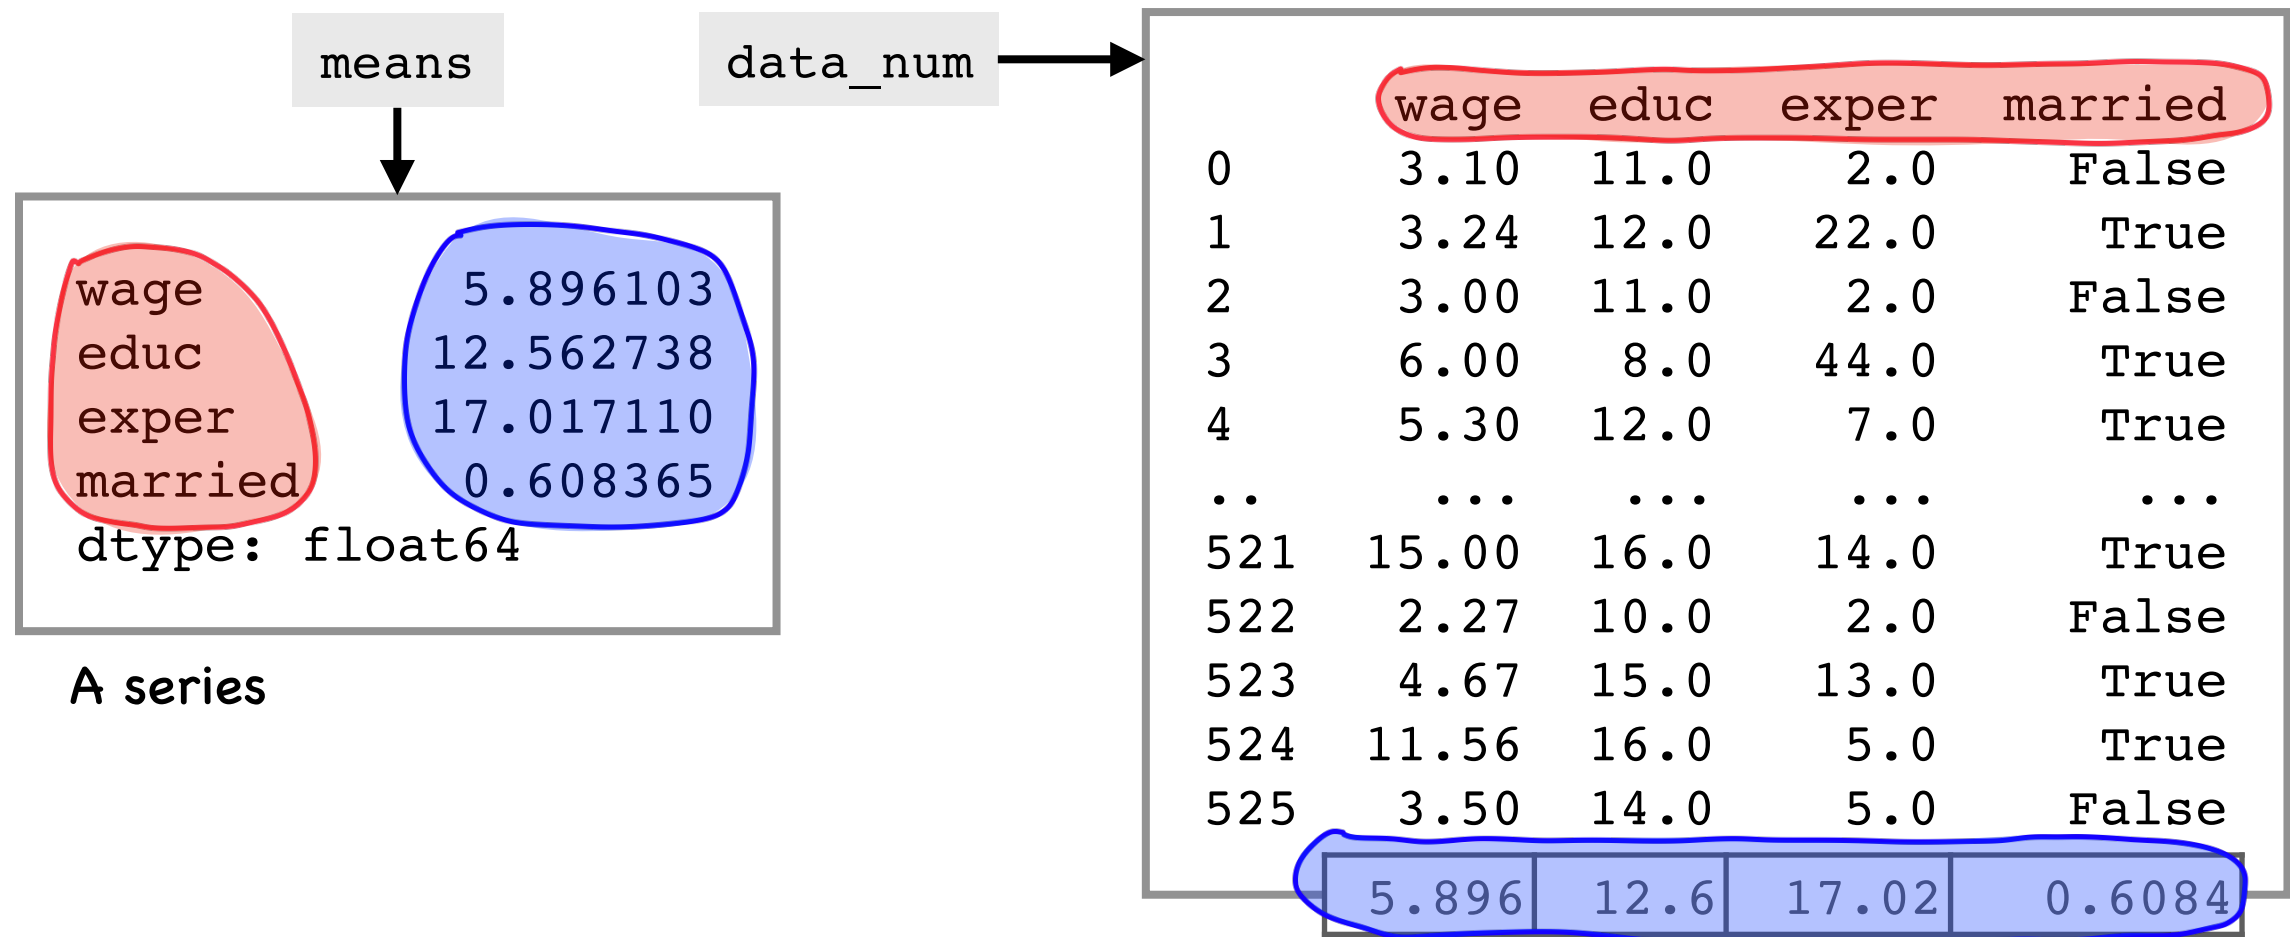
- ✓ Mode



Basics of Descriptive Analytics

- Descriptive measures
 - Measures of centers

`means = data_num.mean()` → Mean values



Basics of Descriptive Analytics

- Descriptive measures
 - Measures of centers

```
means = data_num.mean()
```

means

wage	5.896103
educ	12.562738
exper	17.017110
married	0.608365
dtype:	float64

Proportion of married cases

$$\hat{p} = \frac{1}{n} \sum_{i=1}^n x_i$$

Total number of observations

Total number of married cases

$$x_i = \begin{cases} 1, & \text{if True} \\ 0, & \text{if False} \end{cases}$$

Basics of Descriptive Analytics

- Descriptive measures
 - Measures of centers

```
medians = data_num.median()
```

Median values

means

data_num

wage	4.65
educ	12.00
exper	13.50
married	1.00
dtype: float64	

	wage	educ	exper	married
0	3.10	11.0	2.0	False
1	3.24	12.0	22.0	True
2	3.00	11.0	2.0	False
3	6.00	8.0	44.0	True
4	5.30	12.0	7.0	True
..
521	15.00	16.0	14.0	True
522	2.27	10.0	2.0	False
523	4.67	15.0	13.0	True
524	11.56	16.0	5.0	True
525	3.50	14.0	5.0	False

4.65	12.00	13.50	1.00
------	-------	-------	------

Basics of Descriptive Analytics

- Descriptive measures
 - Measures of variations
 - ✓ Variance
 - ✓ Standard deviation

```
data_num.var()
```

```
wage      13.638884
educ       7.667485
exper    184.203516
married    0.238711
dtype: float64
```

```
data_num.std()
```

```
wage      3.693086
educ       2.769022
exper    13.572160
married    0.488580
dtype: float64
```


Basics of Descriptive Analytics

- Descriptive measures
 - Extreme points
 - ✓ Minimum value
 - ✓ Maximum value

```
data.min()
```

```
wage      0.53
educ       0.0
exper      1.0
gender     F
married    False
dtype: object
```

```
data.max()
```

```
wage      24.98
educ      18.0
exper     51.0
gender     M
married    True
dtype: object
```

Basics of Descriptive Analytics

- Descriptive measures
 - Counts of categorical values

```
data[ 'gender' ].value_counts()
```

M	274
F	252

→ Counts (frequencies) of all unique values
Name: gender, dtype: int64

```
data[ 'gender' ].value_counts(normalize=True)
```

M	0.520913
F	0.479087

→ Proportions of all unique values
Name: gender, dtype: float64

Basics of Descriptive Analytics

- Descriptive measures
 - The `corr()` and `cov()` methods

```
data.corr()
```

	wage	educ	exper	married
wage	1.000000	0.405903	0.112903	0.228817
educ	0.405903	1.000000	-0.299542	0.068881
exper	0.112903	-0.299542	1.000000	0.316984
married	0.228817	0.068881	0.316984	1.000000

```
data.cov()
```

	wage	educ	exper	married
wage	13.638884	4.150864	5.659076	0.412871
educ	4.150864	7.667485	-11.257266	0.093188
exper	5.659076	-11.257266	184.203516	2.101952
married	0.412871	0.093188	2.101952	0.238711

Basics of Descriptive Analytics

- Descriptive measures
 - The `describe()` method

```
wage_summary = data.describe()
```

data_num

	wage	educ	exper
count	526.000000	526.000000	526.000000
mean	5.896103	12.562738	17.01711
std	3.693086	2.769022	13.57216
min	0.530000	0.000000	1.00000
25%	3.330000	12.000000	5.00000
50%	4.650000	12.000000	13.50000
75%	6.880000	14.000000	26.00000
max	24.980000	18.000000	51.00000

The 1st quartile (Q1)

The 2nd quartile (Q2)

The 3rd quartile (Q3)

A data frame

Basics of Descriptive Analytics


- Descriptive measures
 - The `describe()` method

```
wage_summary = data.describe()
```


The 1st quartile (Q1)

The 2nd quartile (Q2)


The 3rd quartile (Q3)




(a) Uniform



(b) Bell shaped



(c) Right skewed




(d) Left skewed

Basics of Descriptive Analytics


- Visualizing data

- The boxplot: five-value summary



Basics of Descriptive Analytics


- Visualizing data
 - The boxplot: five-value summary



Basics of Descriptive Analytics

- Visualizing data
 - The boxplot: five-value summary

```
plt.boxplot(data['wage'],  
            vert=False)  
plt.xlabel('Hourly wages ($)',  
           fontsize=14)  
plt.show()
```




Basics of Descriptive Analytics

- Visualizing data

- Histogram

```
plt.hist(data['wage'], bins=20,  
         color='b', alpha=0.3)  
plt.xlabel('Hourly wages ($)',  
          fontsize=14)  
plt.ylabel('Frequency',  
          fontsize=14)  
plt.show()
```

Number of bins




Basics of Descriptive Analytics

- Visualizing data
 - Scatterplots

```
plt.scatter(data['educ'], data['wage'], color='b', alpha=0.4)

plt.xlabel('Education (Years)', fontsize=14)
plt.ylabel('Hourly Salary (Dollars)', fontsize=14)
plt.show()
```




Pandas for Data Analysis

- Visualizing data

Question 2:

- Visualize how workers' potential experiences affect their hourly wages.

```
plt.scatter(data['exper'], data['wage'], alpha=0.5, color='b')  
plt.xlabel('Potential Experience (Years)', fontsize=14)  
plt.ylabel('Hourly Salary (Dollars)', fontsize=14)  
plt.show()
```




Pandas for Data Analysis

- Visualizing data

Question 2:

- Visualize how married male workers' education years affect their hourly wages.

```
subset = data.loc[data['married'] & (data['gender'] == 'M')]  
  
plt.scatter(subset['educ'], subset['wage'], alpha=0.5, color='b')  
plt.xlabel('Education (Years)', fontsize=14)  
plt.ylabel('Hourly Salary (Dollars)', fontsize=14)  
plt.show()
```



Pandas for Data Analysis

- Visualizing data

Question 2:

- Visualize the distribution of female workers' wages.

```
subset = data.loc[data['gender'] == 'F']  
  
plt.hist(subset['wage'], bins=20, alpha=0.5, color='b')  
plt.xlabel('Hourly Salary (Dollars)', fontsize=14)  
plt.ylabel('Frequency', fontsize=14)  
plt.show()
```

